

Utilities

Tuesday, February 07, 2017
1:16 PM

The Teradata Utilities

1. **BTEQ** – The first Teradata query tool. It imports and exports data and should be used when < 500,000 rows. BTEQ is an excellent report writer.
2. **FastLoad** performs Inserts into empty Teradata tables in 64K blocks. There can be NO Secondary Indexes, Triggers, Referential Integrity, or Join Indexes on the table.
3. **MultiLoad** loads to populated Teradata tables in 64K blocks. MultiLoad supports Inserts, Updates, Deletes, and Upserts. There can be NO Unique Secondary Indexes, Triggers, Referential Integrity, or Join Indexes on the table.
4. **TPump** loads into Empty or Populated Tables and is a lot like MultiLoad except TPump does all of the work at the row level and NOT the block level. It is called TPump because you can turn up or down the pump to increase or decrease the load quantities, and it allows Secondary Indexes, Referential Integrity, Triggers, or Join Indexes.
5. **FastExport** – Exports data off of Teradata at the block level. It is much faster than a BTEQ Export and is used for exports when there are more than 500,000 rows export ed.
6. **TPT** – The Teradata Parallel Transport is one scripting language that combines all of the above utilities into one common utility.

Above are the utilities and their specialties.

Block Level Utilities

1. **FastLoad** performs Inserts into empty Teradata tables in 64K blocks.
2. **MultiLoad** loads to populated Teradata tables in 64K blocks.
3. **FastExport** – Exports data off of Teradata at the 64K block level.
 - No Secondary Indexes (FastLoad)
 - No Unique Secondary Indexes (FastLoad and MultiLoad)
 - No Triggers (FastLoad and MultiLoad)
 - No Join Indexes (FastLoad and MultiLoad)
 - No Referential Integrity (FastLoad and MultiLoad)

Only 30 FastLoad and MultiLoad combinations can run at one time.

Only 60 FastExport jobs can run simultaneously!

The block utilities have limits that prevent too many of them running at the same time. This feature is actually controlled by a new DBS Control parameter named MaxLoadAWT which controls AMP Worker Tasks (AWT). When MaxLoadAWT is set to zero, then it is like going back in time to pre-V2R6.0 where only 15 FastLoad, MultiLoad, and FastExport jobs can run, max.

Row Level Utilities


1. **BTEQ** – The first Teradata query tool. It imports and exports data and should be used when < 500,000 rows. BTEQ is an excellent report writer.
2. **TPump** loads into Empty or Populated Tables and is a lot like MultiLoad except TPump does all of the work at the row level and NOT the block level. It is called TPump because you can turn up or down the pump to increase or decrease the load quantities.

Secondary Indexes allowed
Triggers allowed
Join Indexes allowed
Referential Integrity allowed

You can have an unlimited amount of TPump and BTEQ utilities running simultaneously!

Row Level Utilities are not as fast as block level utilities, but they are really flexible and easy to implement. There are almost no restrictions.

Fast Path Inserts Using Insert/Select

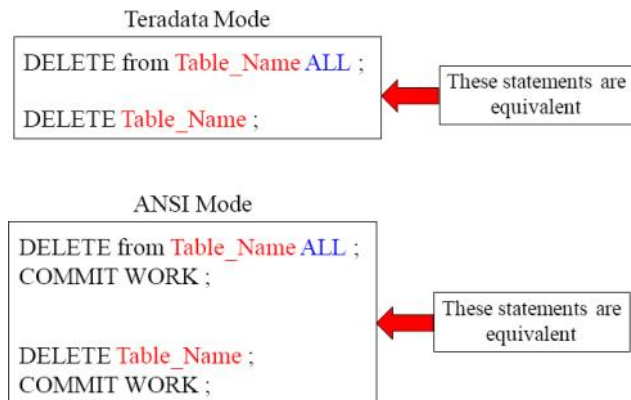


Why is the Fast Path so fast?

- Only one write to the Transient Journal (Now called WAL Journal) because it is easy to Rollback an empty table.
- The Data is copied and written to disk in blocks.
- No Data is redistributed over the BYNET because the tables have the same Primary Index.

Using the Fast Path for Inserts can be really fast.

Fast Path Deletes



- ♦ Data is deleted at the Cylinder Index and Master Index level.
- ♦ ANSI syntax requires the COMMIT WORK statement.

Using the Fast Path for Deletes can be really fast.


Freespace Percent

These **Don't** Honor Freespace Percent

- ♦ SQL Inserts and Updates
- ♦ TPump
- ♦ MultiLoad Insert and Updates to populated tables.

These **DO** Honor Freespace Percent

- ♦ FastLoad
- ♦ Insert/Select into an Empty Table
- ♦ MultiLoad into an Empty Table
- ♦ SQL that adds Fallback to a Table
- ♦ SQL that Creates a Secondary Index



Freespace Percent commands **Teradata** not to completely fill up a cylinder. This is actually set in a DBS control parameter called FreeSpacePercent.

Pasted from <[http://techbus.safaribooksonline.com/book/databases/9781940540191/chapter-1-teradata-utilities-introduction/chap1_sub1_xhtml?query=\(\(Teradata+Load+Utilities\)\)#X2ludGVybmlFsX0h0bWxWaWV3P3htbGliPTk3ODE5NDA1NDAxOTEIMkZiaGFwMV9zdWI2X3hodG1sJnF1ZXJ5PSgoVG9yYWRhdGEIMjBmb2FkTlVlVXRpbGloaWVzKSk=>](http://techbus.safaribooksonline.com/book/databases/9781940540191/chapter-1-teradata-utilities-introduction/chap1_sub1_xhtml?query=((Teradata+Load+Utilities))#X2ludGVybmlFsX0h0bWxWaWV3P3htbGliPTk3ODE5NDA1NDAxOTEIMkZiaGFwMV9zdWI2X3hodG1sJnF1ZXJ5PSgoVG9yYWRhdGEIMjBmb2FkTlVlVXRpbGloaWVzKSk=>)>